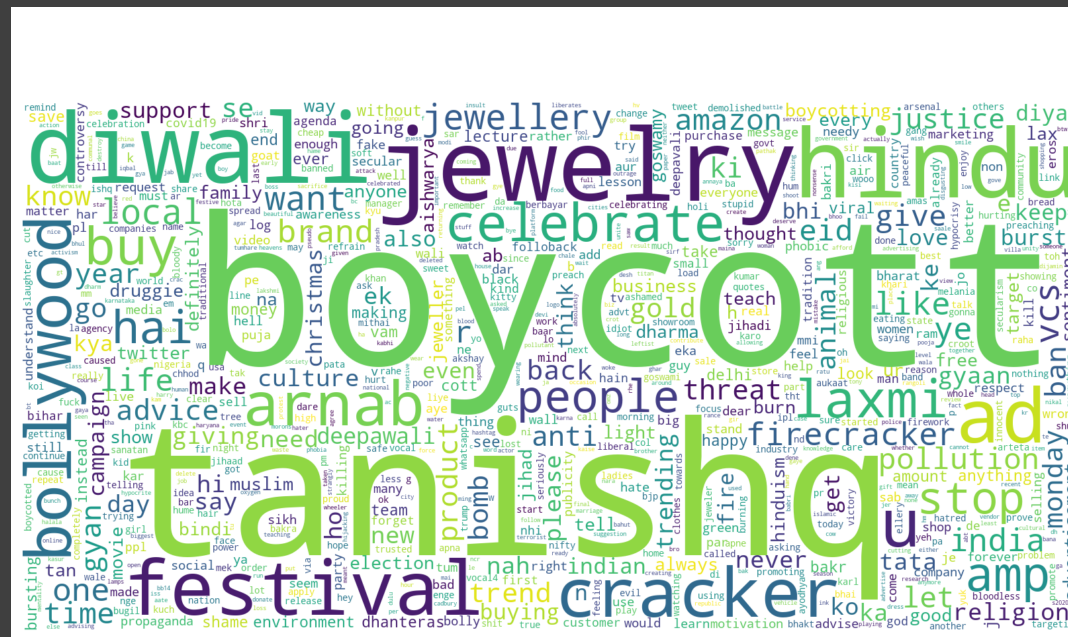


Sentiment Analysis of Tweets about Tanishq Ad Fiasco

Osman Mamun

Data Scientist, Pacific Northwest National Laboratory



Motivation

- During Diwali last year, Tanishq, an Indian Jewelry brand, made an ad featuring interfaith marriage which led to a widespread controversy throughout India. As a result, an initiative to boycott Tanishq was carried out on Twitter with the #BoycottTanishq slogan.
- It is essential for company to understand public sentiment to quickly correct their actions to avoid any negative consequences for the business.
- Here, I analyzed the twitter data related to this ad controversy to identify positive and negative twitter sentiment.

Approach

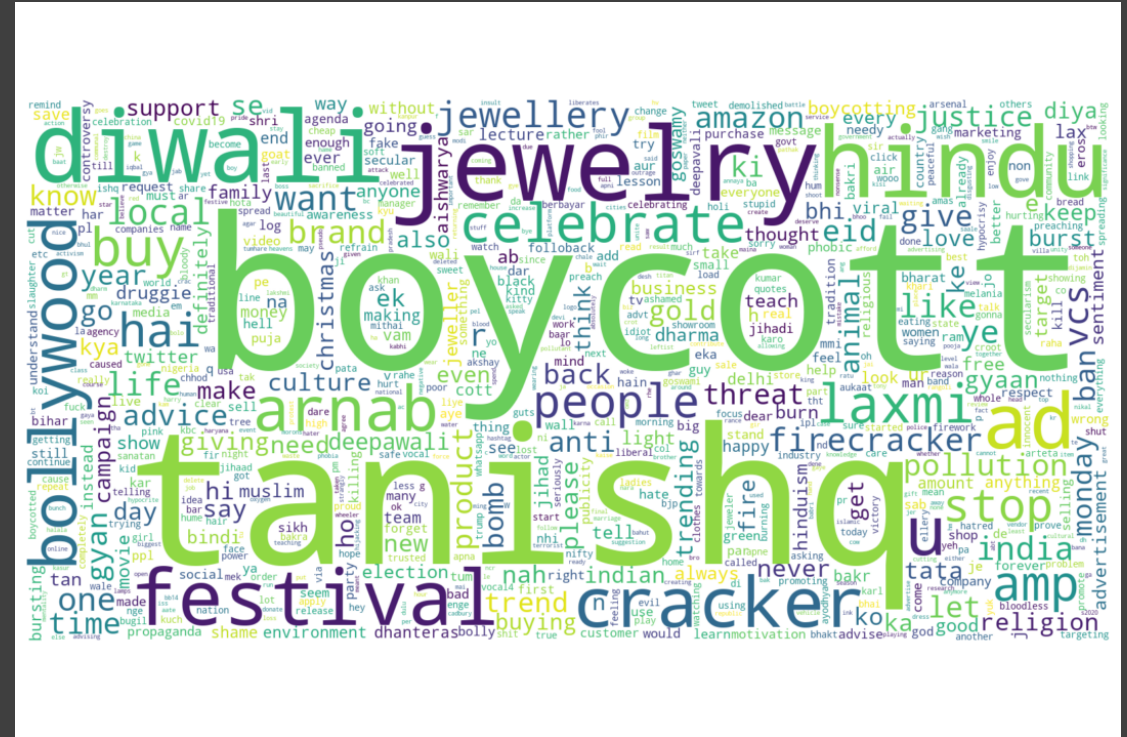
- Collect Twitter Data
- Data Cleaning and Visualization
- Manually Hand Label Data
- Analyze the Feasibility of Pretrained Sentiment Analyzer
- Unsupervised Learning Algorithm
- Supervised Learning Algorithm
- Weak Supervision Model with Pseudo Labeling
- Synthetic Minority Oversampling Technique (SMOTE)

Collect Twitter Data

- Tweepy python package is used to scrape Twitter data with Twitter API
- Sensitive information, i.e., access token, consumer key, secrets etc., are stored in .env file for easy retrieval by the code
- Data can be saved as json file or converted to a pandas dataframe to store selected information.

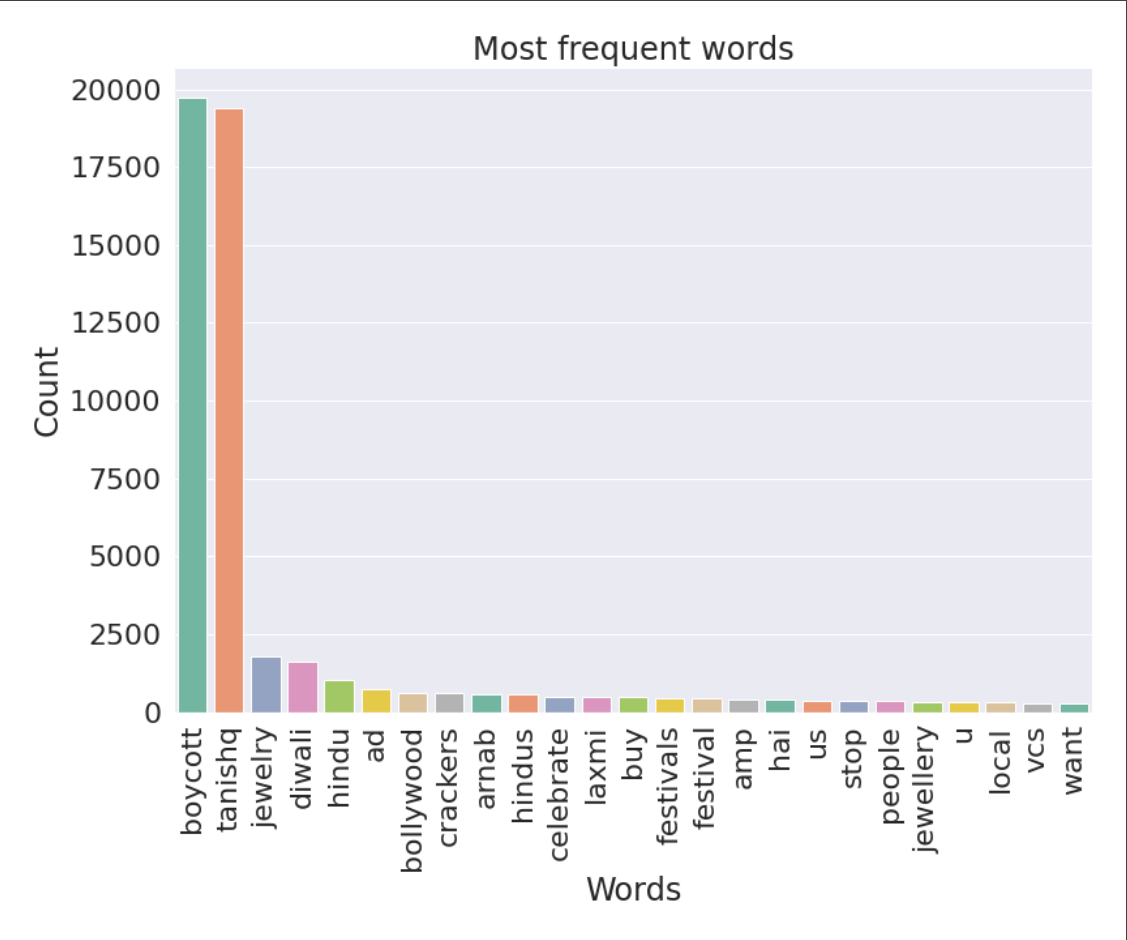
```
1 import os
2 import tweepy
3 from tweepy.parsers import JSONParser
4 import json
5 from dotenv import load_dotenv, find_dotenv
6
7 def fetch_tweets(q=None,
8                 lang=None,
9                 maxResults=None,
10                f_name=None):
11
12     dotenv_path = find_dotenv()
13     load_dotenv(dotenv_path)
14     # Load the authentication key-value pairs into individual variables
15
16     access_token = os.environ.get("access_token")
17     access_token_secret = os.environ.get("access_token_secret")
18     consumer_key = os.environ.get("consumer_key")
19     consumer_secret = os.environ.get("consumer_secret")
20     # Pass OAuth details to tweepy's OAuth handler
21     auth = tweepy.OAuthHandler(consumer_key, consumer_secret)
22     auth.set_access_token(access_token, access_token_secret)
23     api = tweepy.API(auth, parser=JSONParser())
24     json_str = json.dumps(api.search(q=q,
25                                   maxResults=maxResults,
26                                   lang=lang))
27
28     if f_name:
29         with open('../data/raw/' + f_name) as f:
30             json.dump(json_str, f)
31
32     else:
33         return json_str
```

- Hashtags (#), URL, emoji, mentions, and numbers were removed using tweet preprocessing python package.
- All the texts were converted to lowercase and any special characters were removed.
- WordNetLemmatizer and TweetTokenizer from nltk package were used to lemmatize and tokenize the sentences.

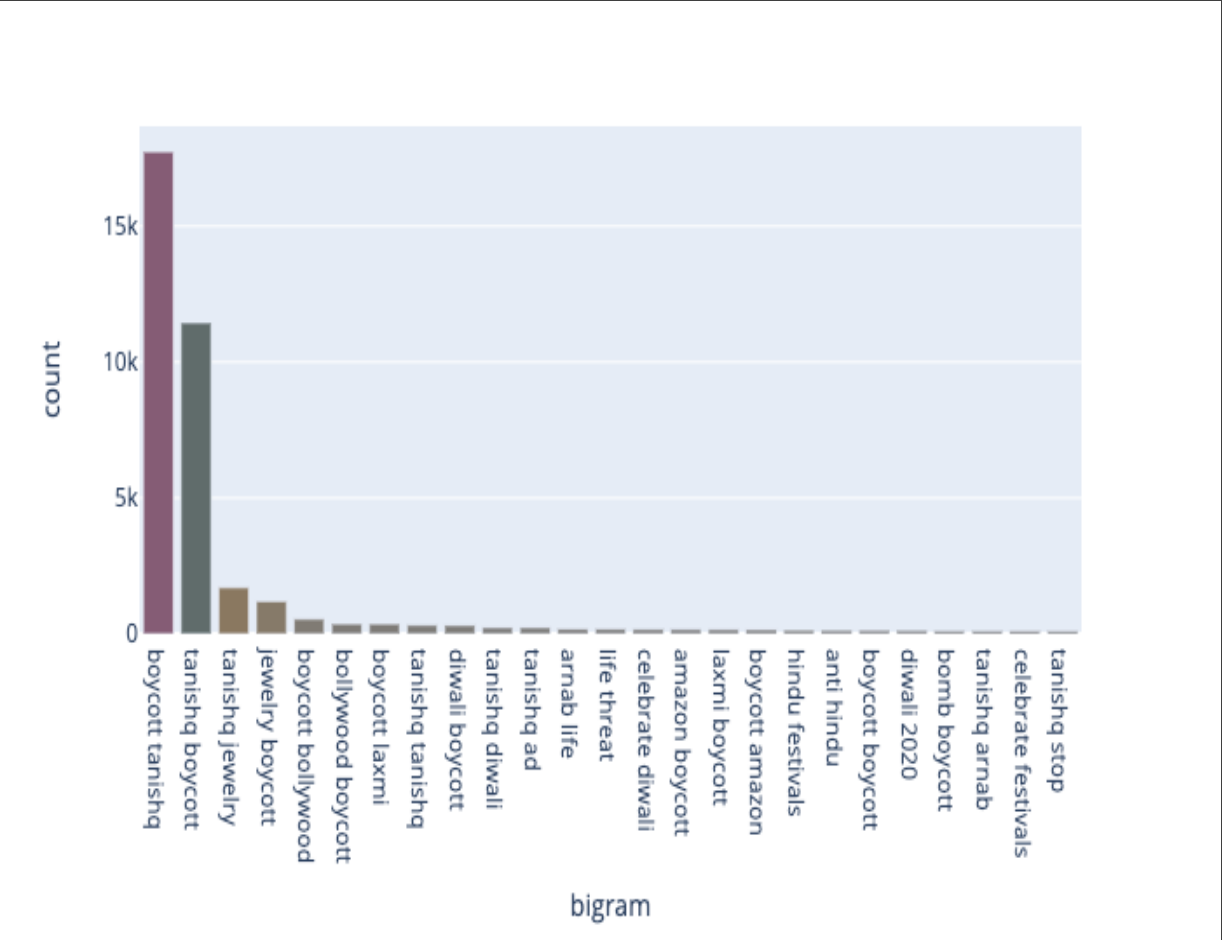


Data Cleaning and Visualization

Word Frequency



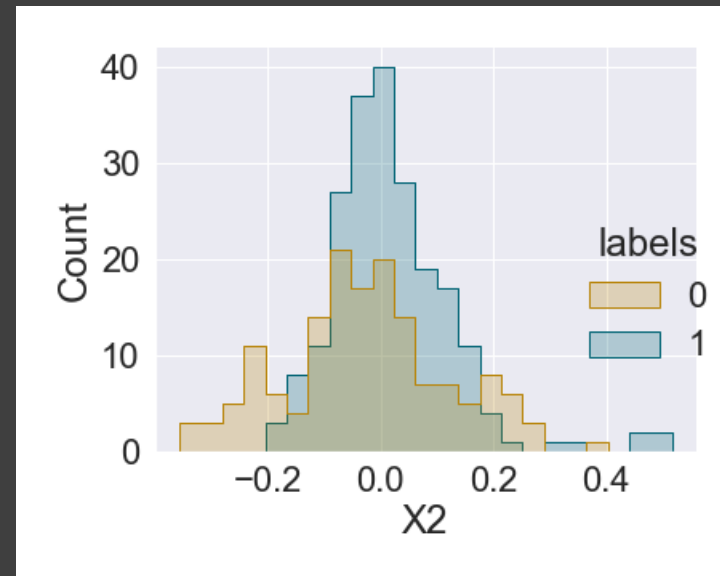
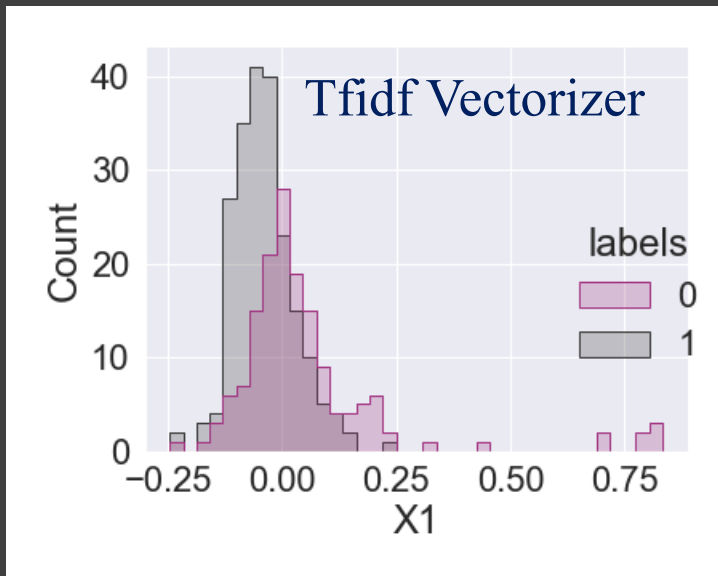
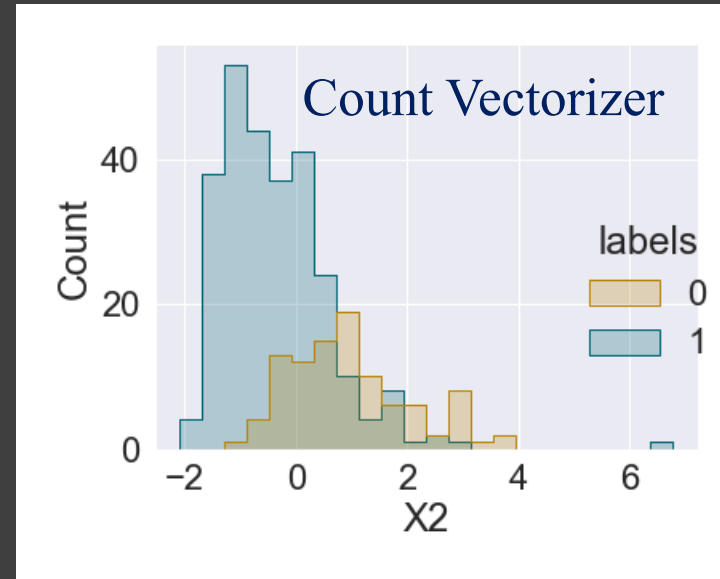
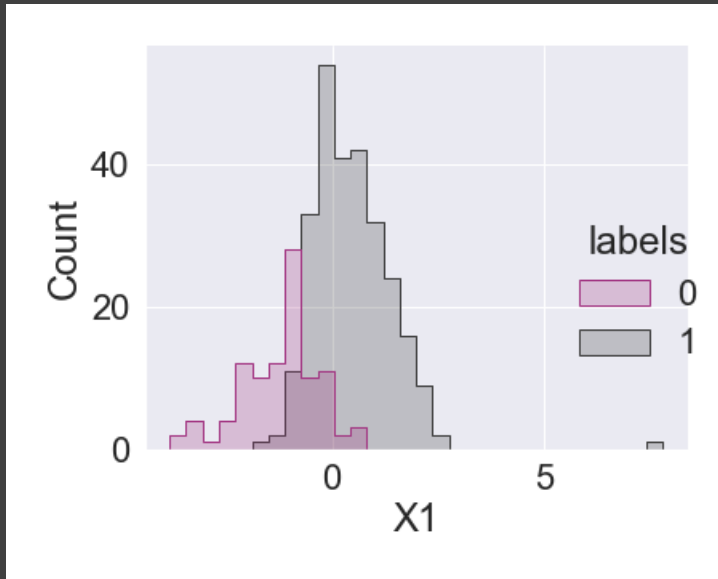
Bigram Frequency



Manually Hand Label Data

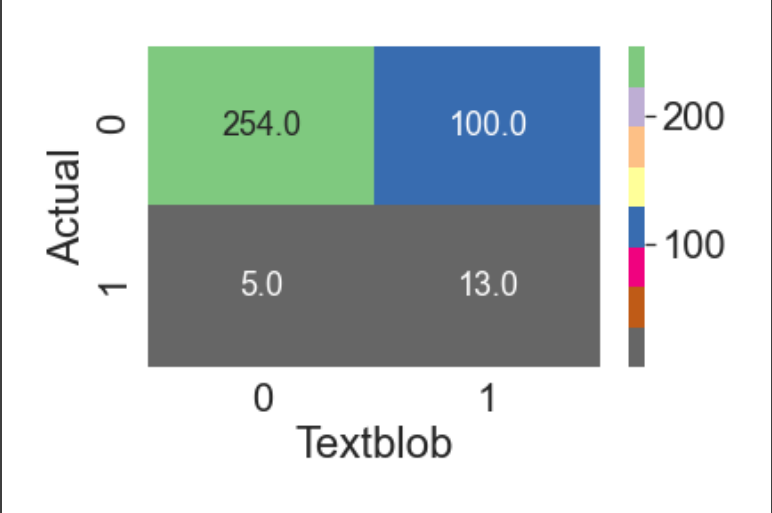
- In order to quantify the performance of different approaches, I manually hand label about 1000 data. It is of note that only about 300 labels were either positive or negative.
- Due to the ambiguity and sarcastic nature of tweets, it was very difficult to label tweets as either positive or negative, even by a human observer.
- For better modeling for business application, it is recommended to manually hand label about 2-3K more data. However, in the interest of time efficiency, we refrained from further hand labeling the data.

Why is It a Challenging Problem



Significance feature
space overlapping with
both the count
vectorizer and tfidf
vectorizer approach.

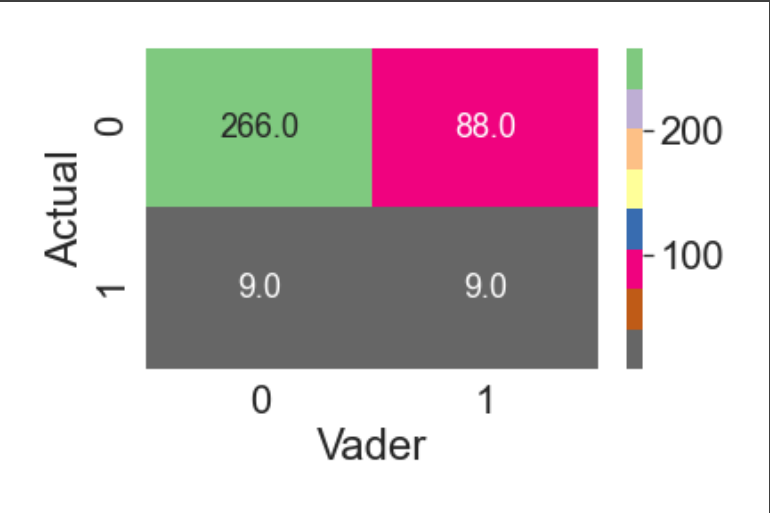
Analyze the Feasibility of Pretrained Sentiment Analyzer



	precision	recall	f1-score	support
0	0.980695	0.717514	0.828711	354
1	0.115044	0.722222	0.198473	18
accuracy	0.717742	0.717742	0.717742	0
macro avg	0.547870	0.719868	0.513592	372
weighted avg	0.938809	0.717742	0.798216	372

ROC AUC: 0.72

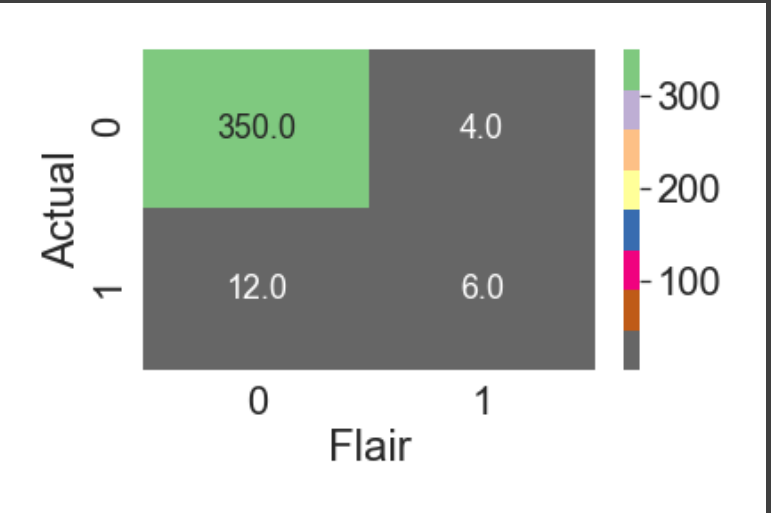
Textblob



	precision	recall	f1-score	support
0	0.967273	0.751412	0.845787	354
1	0.092784	0.500000	0.156522	18
accuracy	0.739247	0.739247	0.739247	0
macro avg	0.530028	0.625706	0.501154	372
weighted avg	0.924959	0.739247	0.812435	372

ROC AUC: 0.63

Vader

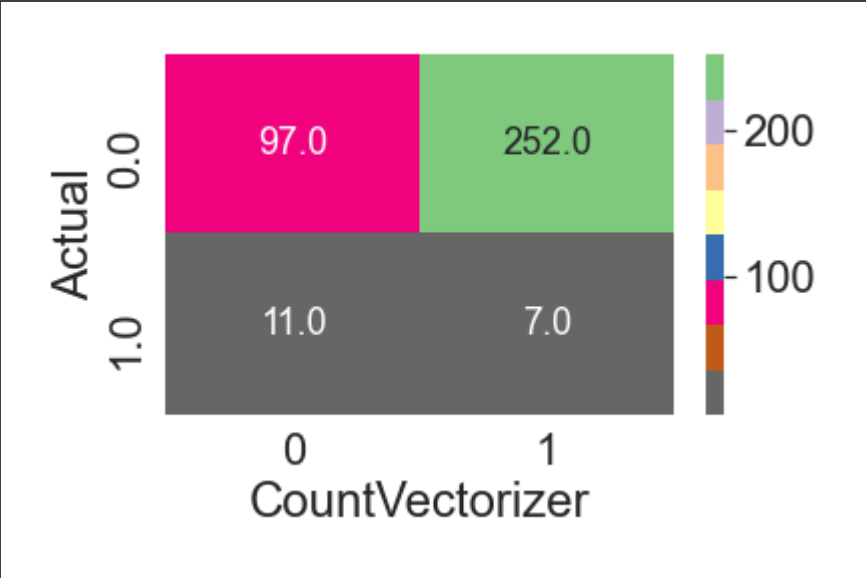


	precision	recall	f1-score	support
0	0.966851	0.988701	0.977654	354
1	0.600000	0.333333	0.428571	18
accuracy	0.956989	0.956989	0.956989	0
macro avg	0.783425	0.661017	0.703113	372
weighted avg	0.949100	0.956989	0.951085	372

ROC AUC: 0.66

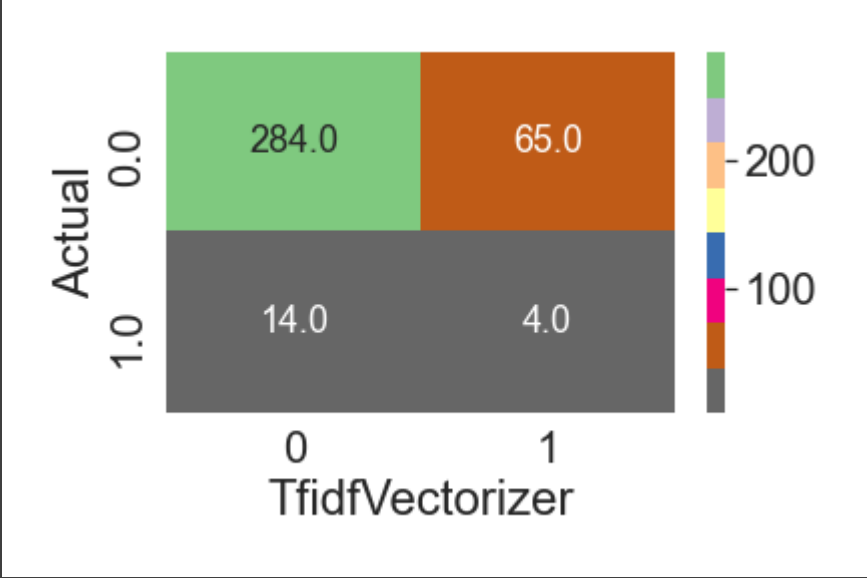
Flair

Unsupervised Learning Algorithm



	precision	recall	f1-score	support
0.0	0.898148	0.277937	0.424508	349
1.0	0.027027	0.388889	0.050542	18
accuracy	0.283379	0.283379	0.283379	0
macro avg	0.462588	0.333413	0.237525	367
weighted avg	0.855423	0.283379	0.406166	367

ROC AUC: 0.33



	precision	recall	f1-score	support
0.0	0.953020	0.813754	0.877898	349
1.0	0.057971	0.222222	0.091954	18
accuracy	0.784741	0.784741	0.784741	0
macro avg	0.505496	0.517988	0.484926	367
weighted avg	0.909121	0.784741	0.839350	367

ROC AUC: 0.51

Pretrained Model
Performance are
inarguably better than
the unsupervised
learning algorithm.

Supervised Learning Algorithm (Count Vectorizer)

	Precision_0	Precision_1	Recall_0	Recall_1	F1_0	F1_1	ROC AUC
Model							
LinearDiscriminantAnalysis	0.98	0.12	0.68	0.80	0.80	0.22	0.74
AdaBoostClassifier	0.97	0.67	0.99	0.40	0.98	0.50	0.69
Perceptron	0.97	0.33	0.95	0.40	0.96	0.36	0.68
XGBClassifier	0.96	1.00	1.00	0.20	0.98	0.33	0.60
LGBMClassifier	0.96	0.50	0.99	0.20	0.97	0.29	0.59

- Very high precision and recall for negative sentiment with most of the classifier
- ROC AUC is also reasonably high
- Low precision and recall for positive sentiment class; however, XGBoost has high precision and Linear Discriminant Analysis has high recall (need proper cross-validation)

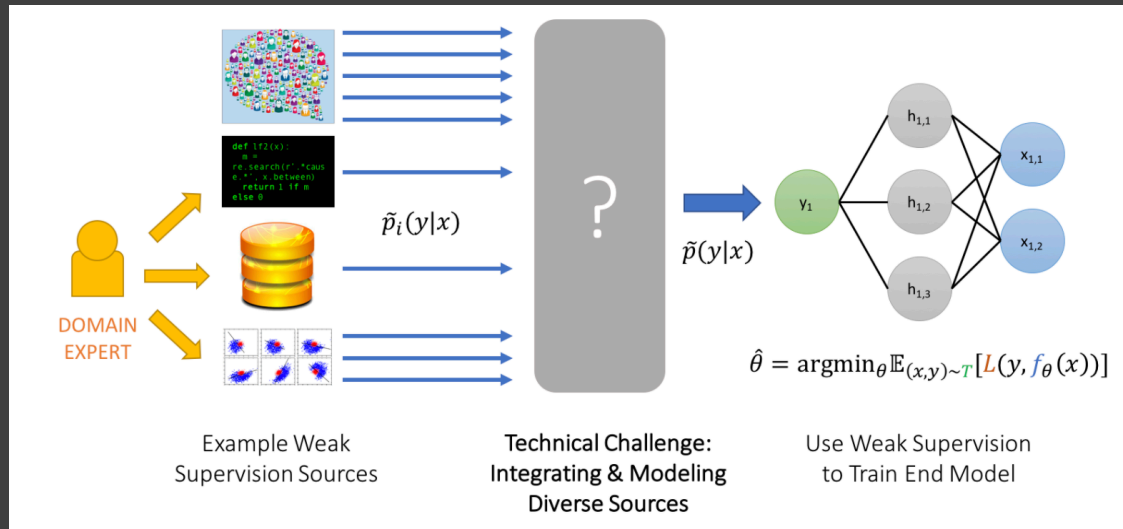
Supervised Learning Algorithm (Tfidf Vectorizer)

	Precision_0	Precision_1	Recall_0	Recall_1	F1_0	F1_1	ROC AUC
Model							
AdaBoostClassifier	0.99	0.80	0.99	0.80	0.99	0.80	0.89
PassiveAggressiveClassifier	0.98	0.38	0.94	0.60	0.96	0.46	0.77
Perceptron	0.98	0.27	0.91	0.60	0.94	0.37	0.75
LinearDiscriminantAnalysis	1.00	0.10	0.46	1.00	0.63	0.18	0.73
DecisionTreeClassifier	0.97	0.67	0.99	0.40	0.98	0.50	0.69

- Overall better performance than the count vectorizer
- ROC AUC is also very impressive across the board
- Low precision and recall for positive sentiment class (high precision in AdaBoost and high recall in AdaBoost and Linear Discriminant).

Weak Supervised Model with Pseudo Labeling

- Snorkel python package was used which is a probabilistic approach to infer the data label given approximate labels
- Requires expensive MCMC sampling
- Once the generative model is trained and labels are generated, a discriminator based on traditional ML algorithm can be trained to decouple expensive generator model



Labeling Function

- Label based on presence and absence of certain keyword or phrases in the tweet
- Learned labels from the previously discussed supervised model.



The rationale for using supervised model learned labels is that a team from Google and CMU reported significant improvement of performance with this approach

Self-training with Noisy Student improves ImageNet classification

Qizhe Xie^{*1}, Minh-Thang Luong¹, Eduard Hovy², Quoc V. Le¹

¹Google Research, Brain Team, ²Carnegie Mellon University

{qizhex, thangluong, qvl}@google.com, hovy@cmu.edu

Weak Supervised Model with Pseudo Labeling

	Precision_0	Precision_1	Recall_0	Recall_1	F1_0	F1_1	ROC	AUC
Model								
LGBMClassifier	1.00	0.10	0.48	1.00	0.65	0.18	0.74	
XGBClassifier	1.00	0.09	0.45	1.00	0.62	0.17	0.72	
BaggingClassifier	1.00	0.08	0.37	1.00	0.54	0.15	0.68	
DecisionTreeClassifier	0.98	0.09	0.51	0.80	0.67	0.15	0.65	
QuadraticDiscriminantAnalysis	0.96	0.12	0.84	0.40	0.90	0.19	0.62	

- With weak supervision, precision for negative sentiment and recall for positive sentiment is very high, but recall for negative and precision for positive is very low
- The algorithm is having trouble navigating a wide overlapping feature space

Synthetic Minority Oversampling Technique (SMOTE)

- Based on the nearest neighbor and efficient interpolation, synthetic minority data samples are generated
- In contrast, ENN (edited nearest neighbor) artificially removes samples of majority class, resulting in oversampling.

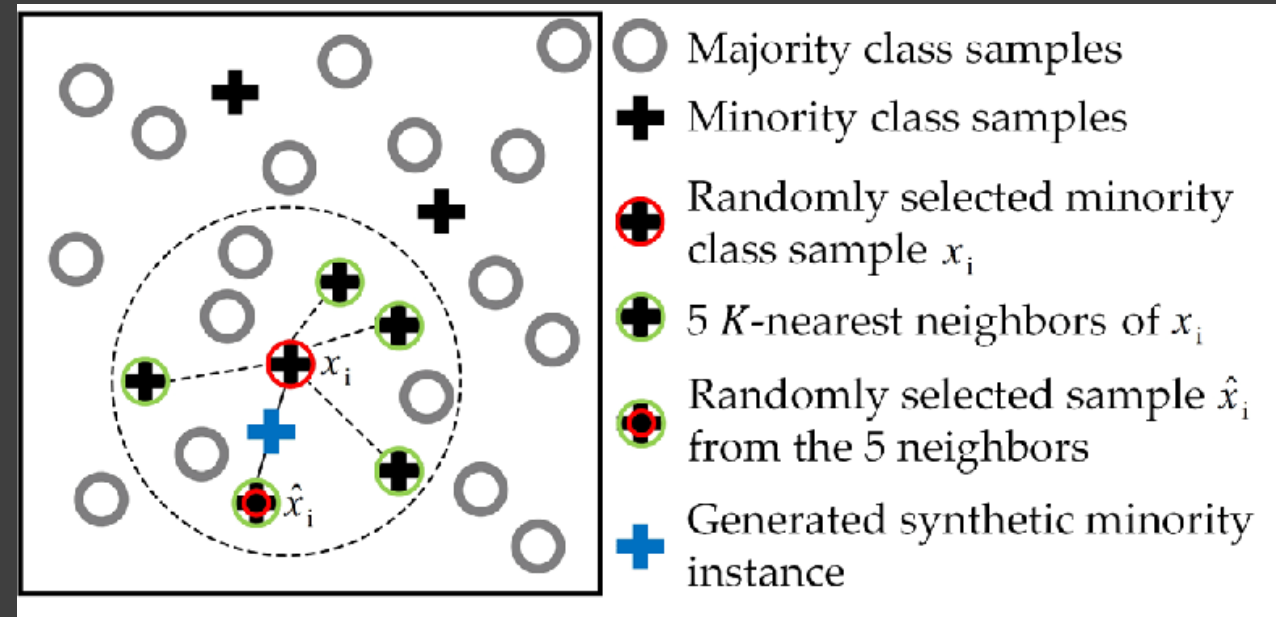
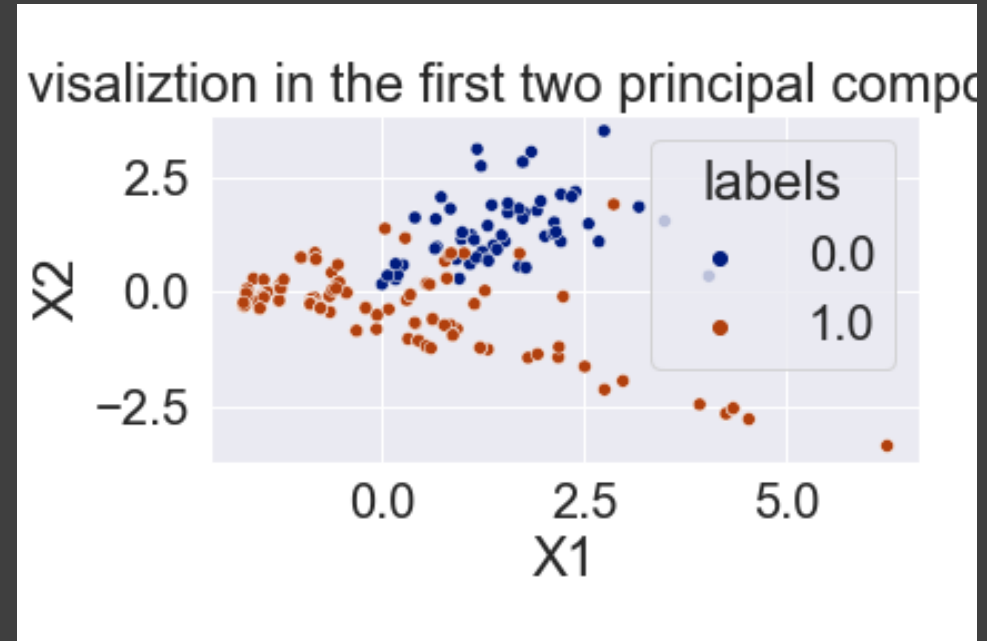
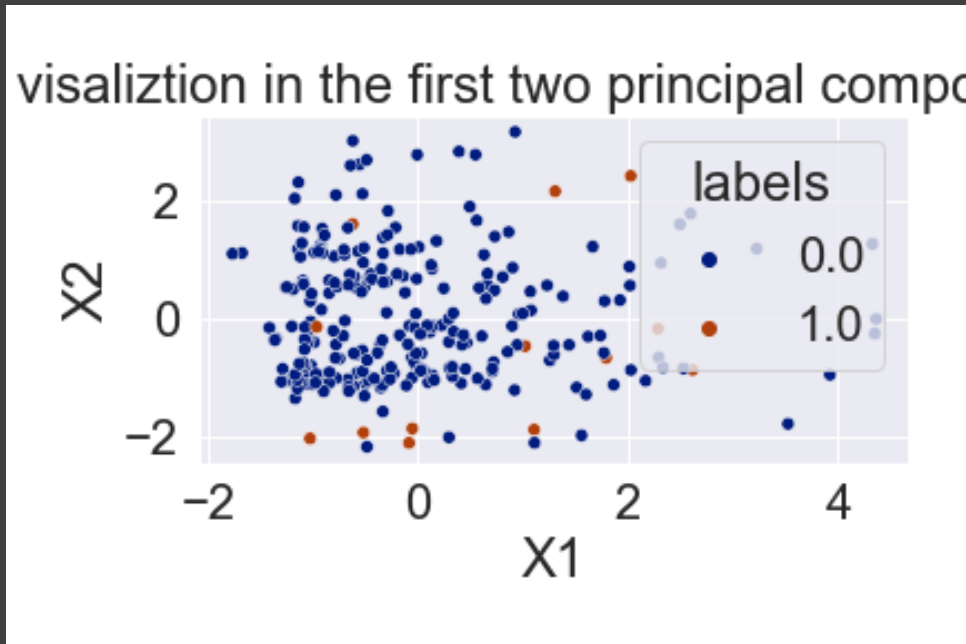


Image credit: <https://www.mdpi.com/2072-4292/11/7/846>

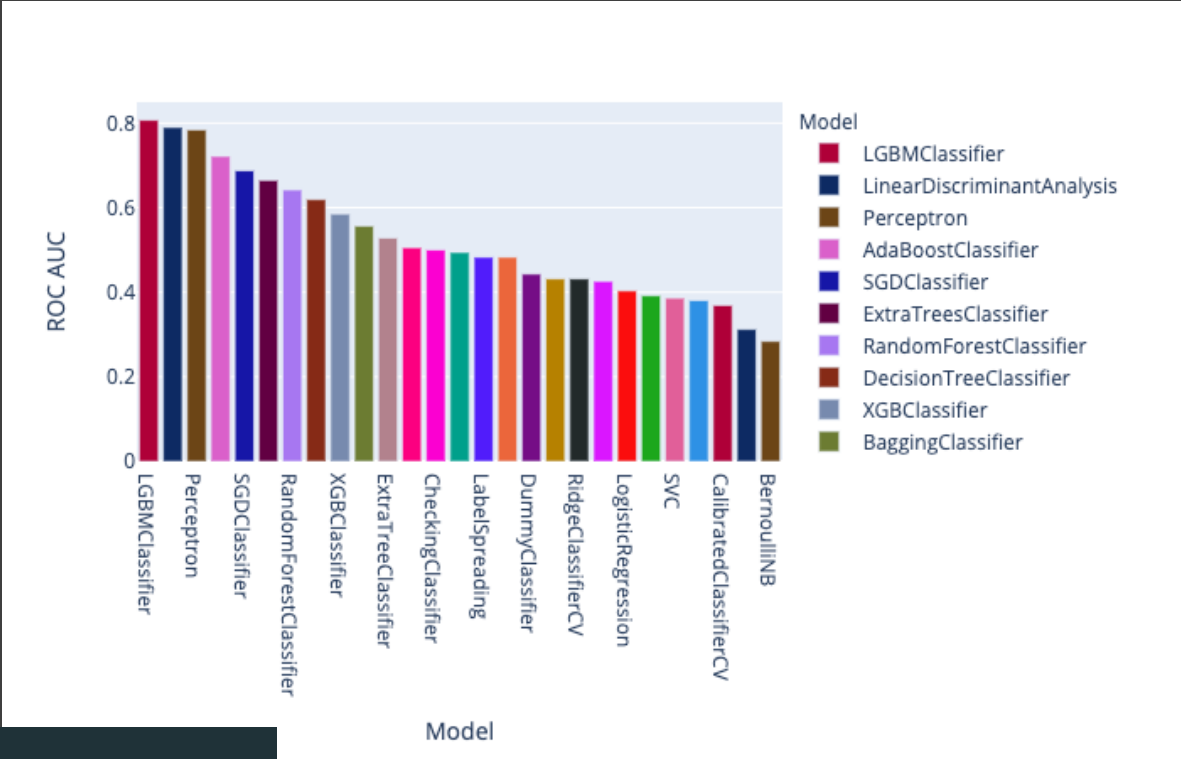
Synthetic Minority Oversampling Technique (SMOTE)



With SMOTEENN, two well defined cluster has been formed by over sampling the minority class and under sampling the majority class in the low PCA region

Synthetic Minority Oversampling Technique (SMOTE)

The performance is not much better than the supervised learning on the labelled dataset only. Next, we will try to perform hyperparameter optimization to improve the performance of SMOTEENN.



	Precision_0	Precision_1	Recall_0	Recall_1	F1_0	F1_1	ROC AUC
Model							
DecisionTreeClassifier	1.00	0.15	0.67	1.00	0.80	0.26	0.83
ExtraTreeClassifier	0.98	0.15	0.75	0.80	0.85	0.26	0.77
LinearDiscriminantAnalysis	1.00	0.11	0.52	1.00	0.68	0.19	0.76
LogisticRegression	1.00	0.10	0.47	1.00	0.64	0.18	0.74
BaggingClassifier	0.98	0.12	0.66	0.80	0.79	0.21	0.73

Hyperparameter Optimization

- AdaBoostClassifier (GridSearchCV): learning_rate, n_estimators , base_estimator
- DecisionTreeClassifier (GridSearchCV): max_depth, min_samples_split, max_features
- XGBClassifier (Bayesian optimization): n_estimators, max_depth, reg_alpha, reg_lambda, eta, gamma, subsample, colsample_bytree

Hyperparameter optimization didn't improve performance much. It seems like, I need to invest more time into tuning the supervised learning hyperparameter to improve the performance

Conclusion

- Out of all the approach, Supervised learning with just the labelled data performed better than other approaches
- For real data analytics, investing more time into labeling data will be crucial to make more reliable model

Acknowledgement: Prasad Seemakurthi (Data Scientist, Navi Ltd.)

Thank You!